

# Exercise Python Code

Ying Guo

## 1 Import all the libraries necessary for this project

```
In [ ]: # data analysis and wrangling
import pandas as pd
import numpy as np
from scipy import stats

#Virtualization
import matplotlib.pyplot as plt
import seaborn as sns

# machine learning
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.feature_selection import RFE
from sklearn.metrics import roc_auc_score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, precision_score
from sklearn.linear_model import LogisticRegression
from sklearn.decomposition import PCA
from imblearn.pipeline import Pipeline
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier

# For oversampling Library (Dealing with Imbalanced Datasets)
from imblearn.over_sampling import SMOTE
```

## 2 Import data

```
In [3]: # Read the file into a DataFrame: df
df = pd.read_csv('exercise_04_train.csv')
df_test = pd.read_csv('exercise_04_test.csv')
```

### 3 Features

In [4]: *# To show all the columns, use the following command*

```
# pd.options.display.max_columns = 4000
```

```
print(df.columns.values)
```

```
print(df.head())
```

```
print(df.shape)
```

```
['x0' 'x1' 'x2' 'x3' 'x4' 'x5' 'x6' 'x7' 'x8' 'x9' 'x10' 'x11' 'x12' 'x13'
 'x14' 'x15' 'x16' 'x17' 'x18' 'x19' 'x20' 'x21' 'x22' 'x23' 'x24' 'x25'
 'x26' 'x27' 'x28' 'x29' 'x30' 'x31' 'x32' 'x33' 'x34' 'x35' 'x36' 'x37'
 'x38' 'x39' 'x40' 'x41' 'x42' 'x43' 'x44' 'x45' 'x46' 'x47' 'x48' 'x49'
 'x50' 'x51' 'x52' 'x53' 'x54' 'x55' 'x56' 'x57' 'x58' 'x59' 'x60' 'x61'
 'x62' 'x63' 'x64' 'x65' 'x66' 'x67' 'x68' 'x69' 'x70' 'x71' 'x72' 'x73'
 'x74' 'x75' 'x76' 'x77' 'x78' 'x79' 'x80' 'x81' 'x82' 'x83' 'x84' 'x85'
 'x86' 'x87' 'x88' 'x89' 'x90' 'x91' 'x92' 'x93' 'x94' 'x95' 'x96' 'x97'
 'x98' 'x99' 'y']
   x0      x1      x2      x3      x4      x5      x6 \
0 -17.933519  6.559220  2.422468 -27.737392 -12.080601 -3.892934  1.067466
1 -37.214754 10.774930  5.404072  21.354738  0.612690 -3.093533  6.161558
2  0.330441 -19.609972 -1.331804 -15.153892 19.710240 19.077300 -1.747110
3 -13.709765 -8.011390 -1.536483  23.129497 27.880879 20.573991 -1.617689
4 -4.202598  7.076210  8.881550 23.600777 26.232164 -14.462320  3.231193

   x7      x8      x9 ...      x91      x92      x93      x94 \
0 0.935953 10.912007 1.107144 ... 11.107047  0.093337  asia  0.421524
1 -0.972156 -5.222169 0.384969 ... -1.991846 15.666187  asia -0.132764
2 0.545570 -1.464609 3.670570 ... 17.132840 -5.333012  asia  1.432308
3 4.129694  1.139928 2.912838 ... 12.292136  4.177925  asia  0.733069
4 -0.069364 -7.310536 -2.268700 ...  6.218743  8.715709  asia -0.977502

   x95      x96      x97      x98      x99 y
0 35.259947  8.994318 -21.000182 -0.686588  2.949106  1
1 -1.192563  3.885024 -37.886523 -7.730392 -1.107330  0
2 -3.435427 -1.133450  7.426099 -5.945534  1.316312  0
3  4.372964 15.529931 29.712153  2.240740  0.477195  0
4 -30.085932 -8.244312 66.540331 -3.478195 -2.869702  1

[5 rows x 101 columns]
(40000, 101)
```

### 4 Exploring features

Categorical features: x34, x35, x68, and x93

Numerical features: x1 - x33, x36 - x67, x69 - x92, and x94 - x99

Features need to be cleaned up: x35, x41, x45, and x68

In [5]: df.describe()

```
Out[5]:
```

	x0	x1	x2	x3	x4 \
count	39989.000000	39989.000000	39993.000000	39991.000000	39992.000000
mean	6.159970	-3.568111	0.223336	-1.742588	0.079437
std	29.098537	17.186748	5.237987	36.601044	21.179065
min	-106.809919	-72.864290	-21.508799	-157.569819	-79.900790
25%	-13.617383	-15.148354	-3.295204	-26.465502	-14.215354
50%	6.247370	-3.660536	0.264994	-1.638876	0.113879
75%	25.570242	7.807474	3.761013	23.044686	14.365631
max	134.592465	71.071223	21.060130	145.566756	89.856546

	x5	x6	x7	x8	x9 \
count	39989.000000	39993.000000	39988.000000	39996.000000	39992.000000
mean	-0.535399	0.015483	-0.011955	-3.055506	-0.023167
std	13.602122	4.110412	2.423051	13.450495	2.472008
min	-55.050043	-15.955862	-9.299563	-54.415601	-9.674058
25%	-9.771613	-2.770450	-1.644516	-12.055884	-1.683043
50%	-0.530463	0.015259	-0.002569	-3.069374	-0.039400
75%	8.673525	2.770460	1.621142	5.910663	1.636558
max	52.628375	18.546313	11.919020	54.262047	9.492780

	...	x90	x91	x92	x94 \
count	...	39995.000000	39995.000000	39992.000000	39990.000000
mean	...	-7.472520	-0.026534	0.016619	-0.000084
std	...	85.885663	9.446348	5.585176	1.135819
min	...	-375.460243	-36.618364	-24.268022	-4.928351
25%	...	-64.312552	-6.390111	-3.764955	-0.771053
50%	...	-5.892459	-0.074239	0.025084	0.001850
75%	...	50.873797	6.360710	3.784911	0.767160
max	...	336.414571	42.835142	23.505468	4.792344

	x95	x96	x97	x98	x99 \
count	39992.000000	39985.000000	39991.000000	39995.000000	39990.000000
mean	0.054600	-0.459762	-4.925135	0.033761	0.120155
std	22.278277	12.702453	34.931541	5.374336	3.116143
min	-101.342320	-57.873114	-140.638773	-22.402508	-13.024105
25%	-14.881499	-8.968785	-28.431741	-3.590052	-1.992603
50%	0.239447	-0.371605	-5.023371	0.031702	0.115059
75%	15.109761	8.128631	18.412348	3.663242	2.230546
max	92.442885	52.159468	147.391902	21.614385	13.208294

	y
count	40000.000000
mean	0.203000
std	0.402238
min	0.000000
25%	0.000000

50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 97 columns]

```
In [6]: # All features with object data type
df.loc[:, df.dtypes == object].head()
```

```
Out[6]:
```

	x34	x35	x41	x45	x68	x93
0	bmw	thur	\$-1306.52	-0.01%	sept.	asia
1	Toyota	wednesday	\$-24.86	0.0%	July	asia
2	bmw	thursday	\$-110.85	0.0%	July	asia
3	Toyota	wed	\$-324.43	0.01%	Apr	asia
4	Toyota	wednesday	\$1213.37	-0.01%	Aug	asia

```
In [7]: # Clean up x35, x41, x45, and x68 on both training and test data
# Training data
Xorg = df
Xorg.x35 = Xorg.x35.replace('wed', 'wednesday')
Xorg.x35 = Xorg.x35.replace(['thursday', 'thur'], 'thursday')
Xorg.x35 = Xorg.x35.replace('fri', 'friday')
Xorg['x41'] = Xorg['x41'].str.replace('$', '').astype('float')
print(Xorg['x41'].head())
Xorg['x45'] = Xorg['x45'].str.replace('%', '').astype('float')/100
print(Xorg['x45'].head())
Xorg.x68 = Xorg.x68.replace('January', 'Jan')
Xorg.x68 = Xorg.x68.replace('July', 'Jul')
Xorg.x68 = Xorg.x68.replace('sept.', 'Sep')
Xorg.x68 = Xorg.x68.replace('Dev', 'Dec')
train = Xorg
print(train.shape)

# Test data
Xorg = dftest
Xorg.x35 = Xorg.x35.replace('wed', 'wednesday')
Xorg.x35 = Xorg.x35.replace(['thursday', 'thur'], 'thursday')
Xorg.x35 = Xorg.x35.replace('fri', 'friday')
Xorg['x41'] = Xorg['x41'].str.replace('$', '').astype('float')
print(Xorg['x41'].head())
Xorg['x45'] = Xorg['x45'].str.replace('%', '').astype('float')/100
print(Xorg['x45'].head())
Xorg.x68 = Xorg.x68.replace('January', 'Jan')
Xorg.x68 = Xorg.x68.replace('July', 'Jul')
Xorg.x68 = Xorg.x68.replace('sept.', 'Sep')
Xorg.x68 = Xorg.x68.replace('Dev', 'Dec')
testdata = Xorg
print(testdata.shape)
```

```

0 -1306.52
1 -24.86
2 -110.85
3 -324.43
4 1213.37
Name: x41, dtype: float64
0 -0.0001
1 0.0000
2 0.0000
3 0.0001
4 -0.0001
Name: x45, dtype: float64
(40000, 101)
0 124.72
1 1273.04
2 -1651.19
3 896.05
4 -1710.27
Name: x41, dtype: float64
0 -0.0001
1 -0.0001
2 0.0000
3 0.0001
4 0.0001
Name: x45, dtype: float64
(10000, 100)

```

## 5 Outliers Detection

I applied median-absolute-deviation (MAD) based outlier detection for all numerical features. I used a threshold of 3.5. A data point with Z score whose absolute value larger than 3.5 is labeled as an outlier.

I found 1927 instances with outliers (about 5% of all cases). Without knowing what each feature actually is, it is hard to decide if these outliers are valid data or wrong inputs. I test training data with and without outliers. The results are very similar. Thus I use the dataset with outliers here.

```

In [8]: allnum = train[train.dtypes == float].columns.tolist()
def mad_based_outlier(points, thresh=3.5):
    if len(points.shape) == 1:
        points = points[:,None]
    median = np.nanmedian(points, axis=0)
    diff = np.sum((points - median)**2, axis=-1)
    diff = np.sqrt(diff)
    med_abs_deviation = np.nanmedian(diff)

    modified_z_score = 0.6745 * diff / med_abs_deviation

```

```

    return modified_z_score > thresh

todrop = list()
for i in range(len(allnum.columns)):
    ind = allnum.iloc[:,i][mad_based_outlier(allnum.iloc[:,i],thresh=3.5)]
    todrop = list(set(todrop+ind.index.tolist()))
print(len(todrop))

```

1927

```

In [9]: noout = train.drop(train.index[[todrop]])
        noout.shape

```

Out[9]: (38093, 101)

## 6 Not many missing data in each feature in both training and test data

I will impute the missing data using median in the analysis (more robust to outliers).

```

In [10]: print(train.describe())

```

```

dfest.info()

```

	x0	x1	x2	x3	x4 \
count	39989.000000	39989.000000	39993.000000	39991.000000	39992.000000
mean	6.159970	-3.568111	0.223336	-1.742588	0.079437
std	29.098537	17.186748	5.237987	36.601044	21.179065
min	-106.809919	-72.864290	-21.508799	-157.569819	-79.900790
25%	-13.617383	-15.148354	-3.295204	-26.465502	-14.215354
50%	6.247370	-3.660536	0.264994	-1.638876	0.113879
75%	25.570242	7.807474	3.761013	23.044686	14.365631
max	134.592465	71.071223	21.060130	145.566756	89.856546

	x5	x6	x7	x8	x9 \
count	39989.000000	39993.000000	39988.000000	39996.000000	39992.000000
mean	-0.535399	0.015483	-0.011955	-3.055506	-0.023167
std	13.602122	4.110412	2.423051	13.450495	2.472008
min	-55.050043	-15.955862	-9.299563	-54.415601	-9.674058
25%	-9.771613	-2.770450	-1.644516	-12.055884	-1.683043
50%	-0.530463	0.015259	-0.002569	-3.069374	-0.039400
75%	8.673525	2.770460	1.621142	5.910663	1.636558
max	52.628375	18.546313	11.919020	54.262047	9.492780

	...	x90	x91	x92	x94 \
count	...	39995.000000	39995.000000	39992.000000	39990.000000
mean	...	-7.472520	-0.026534	0.016619	-0.000084

std	...	85.885663	9.446348	5.585176	1.135819
min	...	-375.460243	-36.618364	-24.268022	-4.928351
25%	...	-64.312552	-6.390111	-3.764955	-0.771053
50%	...	-5.892459	-0.074239	0.025084	0.001850
75%	...	50.873797	6.360710	3.784911	0.767160
max	...	336.414571	42.835142	23.505468	4.792344

	x95	x96	x97	x98	x99 \
count	39992.000000	39985.000000	39991.000000	39995.000000	39990.000000
mean	0.054600	-0.459762	-4.925135	0.033761	0.120155
std	22.278277	12.702453	34.931541	5.374336	3.116143
min	-101.342320	-57.873114	-140.638773	-22.402508	-13.024105
25%	-14.881499	-8.968785	-28.431741	-3.590052	-1.992603
50%	0.239447	-0.371605	-5.023371	0.031702	0.115059
75%	15.109761	8.128631	18.412348	3.663242	2.230546
max	92.442885	52.159468	147.391902	21.614385	13.208294

	y
count	40000.000000
mean	0.203000
std	0.402238
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 97 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10000 entries, 0 to 9999

Data columns (total 100 columns):

x0 9997 non-null float64  
x1 9999 non-null float64  
x2 9998 non-null float64  
x3 9996 non-null float64  
x4 10000 non-null float64  
x5 10000 non-null float64  
x6 9996 non-null float64  
x7 9999 non-null float64  
x8 9997 non-null float64  
x9 9999 non-null float64  
x10 9999 non-null float64  
x11 9997 non-null float64  
x12 10000 non-null float64  
x13 9994 non-null float64  
x14 9998 non-null float64  
x15 9997 non-null float64  
x16 9998 non-null float64

x17 9997 non-null float64  
x18 9998 non-null float64  
x19 9998 non-null float64  
x20 9998 non-null float64  
x21 10000 non-null float64  
x22 10000 non-null float64  
x23 9997 non-null float64  
x24 9996 non-null float64  
x25 9997 non-null float64  
x26 9998 non-null float64  
x27 9995 non-null float64  
x28 9997 non-null float64  
x29 9999 non-null float64  
x30 10000 non-null float64  
x31 9997 non-null float64  
x32 9999 non-null float64  
x33 9997 non-null float64  
x34 9999 non-null object  
x35 9998 non-null object  
x36 9997 non-null float64  
x37 9999 non-null float64  
x38 9999 non-null float64  
x39 9998 non-null float64  
x40 10000 non-null float64  
x41 10000 non-null float64  
x42 9998 non-null float64  
x43 10000 non-null float64  
x44 9999 non-null float64  
x45 9998 non-null float64  
x46 9999 non-null float64  
x47 9999 non-null float64  
x48 9994 non-null float64  
x49 9997 non-null float64  
x50 10000 non-null float64  
x51 9998 non-null float64  
x52 9999 non-null float64  
x53 9999 non-null float64  
x54 10000 non-null float64  
x55 9999 non-null float64  
x56 9999 non-null float64  
x57 9997 non-null float64  
x58 9998 non-null float64  
x59 9997 non-null float64  
x60 9997 non-null float64  
x61 9997 non-null float64  
x62 9996 non-null float64  
x63 9998 non-null float64  
x64 10000 non-null float64



```
x65    9998 non-null float64
x66    9998 non-null float64
x67    9996 non-null float64
x68    10000 non-null object
x69    9997 non-null float64
x70    9998 non-null float64
x71    9999 non-null float64
x72    10000 non-null float64
x73    9995 non-null float64
x74    9997 non-null float64
x75    9998 non-null float64
x76    10000 non-null float64
x77    9995 non-null float64
x78    9999 non-null float64
x79    9997 non-null float64
x80    9998 non-null float64
x81    9998 non-null float64
x82    9997 non-null float64
x83    9999 non-null float64
x84    10000 non-null float64
x85    9997 non-null float64
x86    9997 non-null float64
x87    9998 non-null float64
x88    9998 non-null float64
x89    9998 non-null float64
x90    9997 non-null float64
x91    10000 non-null float64
x92    10000 non-null float64
x93    9999 non-null object
x94    9998 non-null float64
x95    9999 non-null float64
x96    9998 non-null float64
x97    9996 non-null float64
x98    9998 non-null float64
x99    9995 non-null float64
dtypes: float64(96), object(4)
memory usage: 7.6+ MB
```

## 7 check number of instances in each category for four categorical features

```
In [11]: print(train.x34.value_counts(dropna = False))
          print(train.x35.value_counts(dropna = False))
          print(train.x68.value_counts(dropna = False))
          print(train.x93.value_counts(dropna = False))
```

```

volkswagon    12557
Toyota        10922
bmw           7288
Honda         5195
tesla         2286
chrystler     1209
nissan         339
ford          160
mercades      26
chevrolet     10
NaN           8
Name: x34, dtype: int64
wednesday    20756
thursday     17726
tuesday       898
friday        550
monday        59
NaN           11
Name: x35, dtype: int64
Jul          11146
Jun           9289
Aug           8115
May           4833
Sep           3441
Apr           1638
Oct            886
Mar            397
Nov            160
Feb            52
Dec            20
Jan            12
NaN            11
Name: x68, dtype: int64
asia         35434
america       3136
euorpe        1423
NaN            7
Name: x93, dtype: int64

```

```

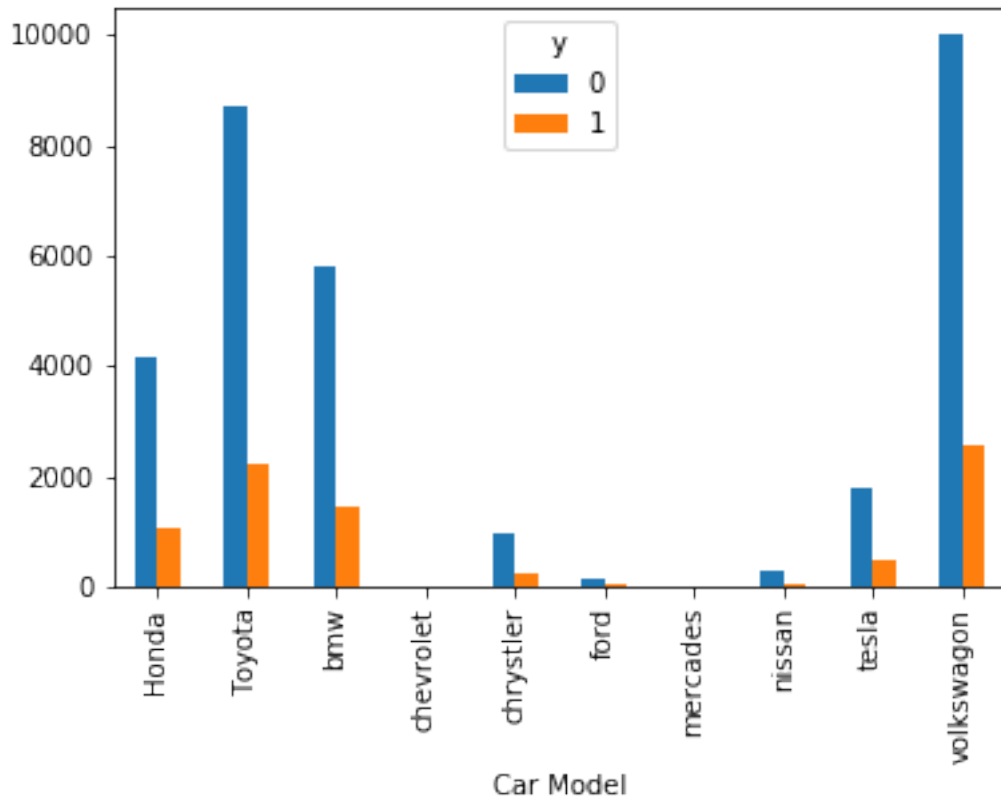
In [12]: %matplotlib inline
         pd.crosstab(train.x34,train.y).plot(kind='bar')
         #plt.title('')
         plt.xlabel('Car Model')
         #plt.ylabel('Type of loans')

```

```


Out[12]: Text(0.5,0,'Car Model')

```



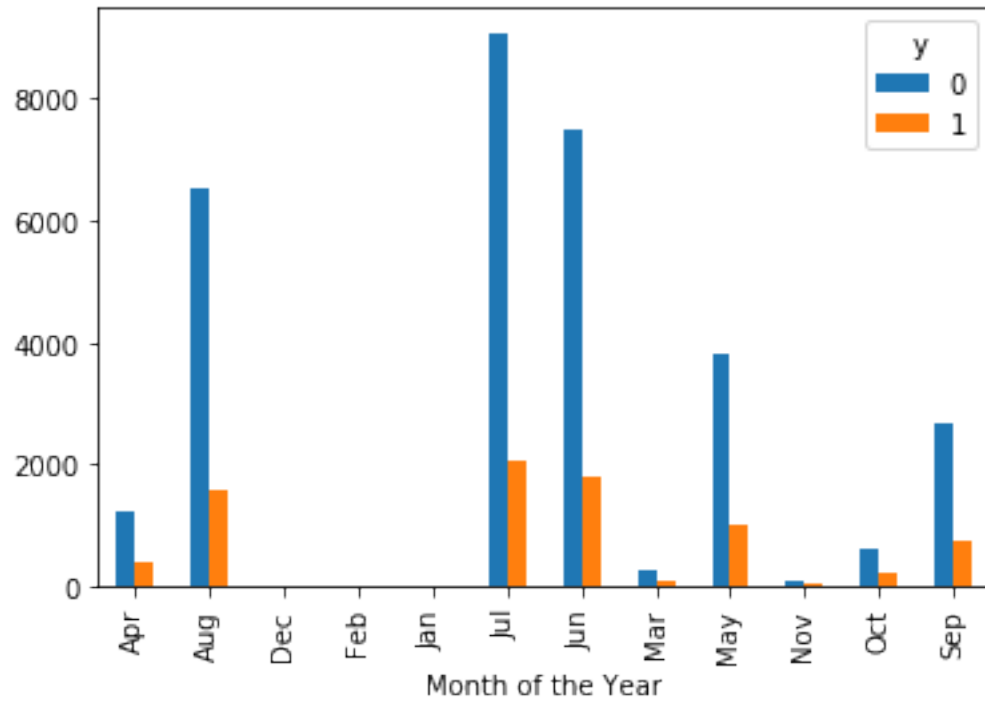
```
In [13]: %matplotlib inline
pd.crosstab(train.x35, train.y).plot(kind='bar')
#plt.title('')
plt.xlabel('Days of the Week')
#plt.ylabel('Type of loans')
```

Out[13]: Text(0.5,0,'Days of the Week')




```
In [14]: %matplotlib inline
pd.crosstab(train.x68,train.y).plot(kind='bar')
#plt.title('')
plt.xlabel('Month of the Year')
#plt.ylabel('Type of loans')
```

```
Out[14]: Text(0.5,0,'Month of the Year')
```



```
In [15]: %matplotlib inline
pd.crosstab(train.x93,train.y).plot(kind='bar')
#plt.title('')
plt.xlabel('Contentient')
#plt.ylabel('Type of loans')
```

```
Out[15]: Text(0.5,0,'Contentient')
```



```
In [16]: train.groupby(['y'], as_index=False).mean()
```

```
Out[16]:
```

	y	x0	x1	x2	x3	x4	x5	x6	\
0	0	7.131102	-4.431030	-0.010624	0.107336	0.164672	0.077979	0.018320	
1	1	2.347929	-0.179797	1.141973	-9.004659	-0.255125	-2.943127	0.004344	

		x7	x8	...	x89	x90	x91	x92	\
0	0	-0.019425	-2.995378	...	0.007417	-7.688762	-0.016384	0.036807	
1	1	0.017378	-3.291546	...	-0.005237	-6.623663	-0.066380	-0.062670	

		x94	x95	x96	x97	x98	x99
0	0	-0.001939	0.054440	0.139299	-8.166448	0.015663	-0.034353
1	1	0.007200	0.055231	-2.812450	7.798984	0.104801	0.726767

[2 rows x 97 columns]

```
In [17]: train[["x34", "y"]].groupby(['x34'], as_index=False).mean()
```

```
Out[17]:
```

	x34	y
0	Honda	0.203465
1	Toyota	0.204633
2	bmw	0.200878
3	chevrolet	0.000000

```

4 chrysler 0.205955
5   ford 0.193750
6   mercades 0.115385
7   nissan 0.194690
8   tesla 0.206474
9   volkswagon 0.202437

```

```
In [18]: train[["x35", "y"]].groupby(['x35'], as_index=False).mean()
```

```

Out[18]:      x35      y
0   friday 0.174545
1   monday 0.423729
2  thursday 0.171782
3   tuesday 0.344098
4  wednesday 0.223694

```

```
In [19]: train[["x68", "y"]].groupby(['x68'], as_index=False).mean()
```

```

Out[19]:      x68      y
0    Apr 0.254579
1    Aug 0.194701
2    Dec 0.400000
3    Feb 0.442308
4    Jan 0.333333
5    Jul 0.186793
6    Jun 0.195177
7    Mar 0.282116
8    May 0.213946
9    Nov 0.337500
10   Oct 0.277652
11   Sep 0.217088

```

```
In [20]: train[["x93", "y"]].groupby(['x93'], as_index=False).mean()
```

```

Out[20]:      x93      y
0  america 0.215242
1    asia 0.202150
2  euorpe 0.196767

```


## 8 Exploring Numerical Features : a few example

```

In [21]: train.groupby('y').x1.plot(kind='kde')
         plt.legend('01', ncol=2, loc='upper left')
         plt.xlabel('x1')


```

```
Out[21]: Text(0.5,0,'x1')
```



```
In [22]: train.groupby('y').x2.plot(kind='kde')
plt.legend('01', ncol=2, loc='upper left')
plt.xlabel('x2')
```


```
Out[22]: Text(0.5,0,'x2')
```





```
In [23]: train.groupby('y').x3.plot(kind='kde')
plt.legend('01', ncol=2, loc='upper left')
plt.xlabel('x3')
```

Out[23]: Text(0.5,0,'x3')



## 9 Correlation among Features

I want to find if there are any highly correlated features. I do not want collinearity in the feature space.

The highest absolute correlation for a pair of features is 0.412, suggesting the feature space does not contain highly correlated features.

```
In [24]: def get_redundant_pairs(df):
    '''Get diagonal and lower triangular pairs of correlation matrix'''
    pairs_to_drop = set()
    cols = df.columns
    for i in range(0, df.shape[1]):
        for j in range(0, i+1):
            pairs_to_drop.add((cols[i], cols[j]))
    return pairs_to_drop
```

```
def get_top_abs_correlations(df, n=5):
    au_corr = df.corr().abs().unstack()
    labels_to_drop = get_redundant_pairs(df)
    au_corr = au_corr.drop(labels=labels_to_drop).sort_values(ascending=False)
    return au_corr[0:n]

print("Top Absolute Correlations")
print(get_top_abs_correlations(allnum, 3))
```

Top Absolute Correlations

x41 x44 0.411508

x90 x95 0.375148

x80 x90 0.373060

dtype: float64

## 10 Divide data into features and target

Here I used all data with and without outliers. The results are similar. Here I used the data with outliers.

In [25]: *# Keep all outliers in the data*

```
x = train.drop("y", axis=1)
```

```
y = train["y"]
```

```
print(x.shape)
```

```
print(y.shape)
```

```
# Remove all outliers in the data
```

```
#x = noout.drop("y", axis=1)
```

```
#y = noout["y"]
```

```
#print(x.shape)
```

```
#print(y.shape)
```

```
(40000, 100)
```

```
(40000,)
```

In [26]: *# Split the original training data into 70% training data and 30% test data*

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, stratify = y,
                                                    random_state = 42)
```

## 11 Building Model 1 : Random Forest with PCA

I built the RF classifier on the reduced dimensions provided by PCA, tuning hyper-parameters with grid search on the number of components for PCA, the number of trees, quality of the split

criteria, number of features to consider when looking for the best split for RF, with 5-fold cross-validation. SMOTE was performed right after PCA. The best performing model selects 110 components from PCA, and uses 500 trees, entropy as split criteria, and the square root of the total number of features as the number of features to consider when looking for the best split for RF. The accuracy score is 0.965, and the ROC AUC score is 0.9824.

In [27]: `from imblearn.pipeline import Pipeline`

```
# We create the preprocessing pipelines for both numeric and categorical data.
numeric_features = train.loc[:, train.dtypes == float].columns.tolist()
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
rf = RandomForestClassifier()
pca = PCA()
clf = Pipeline(steps=[('preprocessor', preprocessor),
    ('pca', pca),
    ('sampling', SMOTE()),
    ('classifier', rf)])

param_grid = { 'pca__n_components': [50,100, 110,124],
    'classifier__n_estimators': [500],
    'classifier__max_features': ['sqrt'],
    'classifier__criterion': ['entropy']
}

rf_cv = GridSearchCV(clf, param_grid, cv = 5)

rf_cv.fit(x_train, y_train.values.ravel())

y_pred_proba = rf_cv.predict_proba(x_test)[: ,1]
# Evaluate test-set roc_auc_score
```

```

rf_roc_auc = roc_auc_score(y_test, y_pred_proba)

print("model score: %.3f" % rf_cv.score(x_test, y_test))
# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(rf_roc_auc))
print("Tuned Random Forest Parameter: {}".format(rf_cv.best_params_))
print("Tuned Random Forest Accuracy: {}".format(rf_cv.best_score_))

```

model score: 0.966

ROC AUC score: 0.9821

Tuned Random Forest Parameter: {'classifier\_\_criterion': 'entropy',  
'classifier\_\_max\_features': 'sqrt', 'classifier\_\_n\_estimators': 500,  
'pca\_\_n\_components': 110}

Tuned Random Forest Accuracy: 0.9617857142857142

In [28]: predictions\_rf = rf\_cv.predict(x\_test)

```

# Classification Report of Prediction
print("Classification Report:")
print(classification_report(y_test, predictions_rf))
# Confusion Matrix for predictions made
conf2 = confusion_matrix(y_test, predictions_rf)
print(conf2)
# Plot Confusion Matrix
label = ["0", "1"]
sns.heatmap(conf2, annot=True, xticklabels=label, yticklabels=label, fmt='g')

```

Classification Report:


	precision	recall	f1-score	support
0	0.97	0.99	0.98	9564
1	0.95	0.88	0.91	2436
micro avg	0.97	0.97	0.97	12000
macro avg	0.96	0.93	0.95	12000
weighted avg	0.97	0.97	0.97	12000

```

[[9442 122]
 [ 289 2147]]

```

Out[29]: <matplotlib.axes.\_subplots.AxesSubplot at 0x10a170b38>



```
In [30]: floatlist = train.loc[:, train.dtypes == float].columns.tolist()
         numcat = rf_cv.best_estimator_.named_steps['preprocessor'].named_transformers_
         bb=numcat['cat'].named_steps["onehot"].get_feature_names()
         feature = np.concatenate((floatlist,bb))
         print(feature.shape)
         #feature
```

(126,)

```
In [31]: pcaname = rf_cv.best_estimator_.named_steps['pca']
         #pcaname.explained_variance_ratio_
```


```
In [32]: a=[];
         for i in range(110):
             a.append( "Comp" + str(i))
         comp = pd.DataFrame(pcaname.components_, columns=feature, index=a)
         c7 = comp.loc['Comp7']#
         ind = c7.abs().sort_values(ascending = False).index.tolist()
         c7.loc[ind].to_csv('out1.csv')
```

```
In [33]: a=[];
         for i in range(22):
             a.append( "Comp" + str(i))
         #type(a)
```

```
In [34]: RF = rf_cv.best_estimator_.named_steps['classifier']
```

```
In [35]: plt.figure(figsize=(110,120))
Importance = pd.DataFrame({'Importance':RF.feature_importances_[0:22]*100},
                           index = a)
Importance.sort_values(
    'Importance', axis=0, ascending=True).plot(kind='barh', color='#875FDB')
plt.xlabel('RF: Variable Importance')
plt.gca().legend_ = None
plt.savefig('RF')
```

<Figure size 7920x8640 with 0 Axes>



```
In [36]: testy_rf_pred_proba = rf_cv.predict_proba(testdata)[: ,1]
pd.DataFrame(testy_rf_pred_proba).to_csv('resultpcar fsmote.csv')
```

## 12 Building Model 2: Gradient Boosting Machine

I used pipelines to build Gradient Boosting Machine. I first created the preprocessing pipelines for both numerical and categorical data. For numerical data, I imputed the missing data with median and standardized the data. For categorical data, I imputed missing value as the most frequent value and created dummy variable (ignoring the missing data). I split the data into 70% training data and 30% test data. I build Gradient Boosting Machine and used grid search on hyperparameter learning rate, maximum depth of the individual estimators, number of features to consider when looking for the best split, with 5-fold cross-validation. The best performing GBM uses 0.2 learning

rate, maximum depth as 8, and the square root of the total number of features as the number of features to consider when looking for the best split. The accuracy score is 0.953, and the ROC AUC score is 0.9834.

In [37]:

```
# We create the preprocessing pipelines for both numeric and categorical data.
numeric_features = train.loc[:, train.dtypes == float].columns.tolist()
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler()))

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
gbm = GradientBoostingClassifier()
clf = Pipeline(steps=[('preprocessor', preprocessor),
    # ('sampling', SMOTE()),
    ('classifier', gbm)])

#clf.fit(X_train, y_train)
param_grid = {
    # "loss": ["deviance"],
    "classifier__learning_rate": [0.01, 0.075, 0.2],
    # "min_samples_split": np.linspace(0.1, 0.5, 12),
    # "min_samples_leaf": np.linspace(0.1, 0.5, 12),
    "classifier__max_depth": [3, 8],
    "classifier__max_features": ["log2", "sqrt"],
    # "criterion": ["friedman_mse", "mae"],
    # "subsample": [0.5, 0.618, 0.8, 0.85, 0.9, 0.95, 1.0],
    # "n_estimators": [10]
}

rf_cv = GridSearchCV(clf, param_grid, cv = 5)

rf_cv.fit(x_train, y_train.values.ravel())
```

```

y_pred_proba = rf_cv.predict_proba(x_test)[:,-1]
# Evaluate test-set roc_auc_score
rf_roc_auc = roc_auc_score(y_test, y_pred_proba)

print("model score: %.3f" % rf_cv.score(x_test, y_test))
# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(rf_roc_auc))
print("Tuned Random Forest Parameter: {}".format(rf_cv.best_params_))
print("Tuned Random Forest Accuracy: {}".format(rf_cv.best_score_))

```

model score: 0.953

ROC AUC score: 0.9834

Tuned GBM Parameter: {'classifier\_\_learning\_rate': 0.2, 'classifier\_\_max\_depth': 8, 'classifier\_\_max\_features': 'sqrt'}

Tuned GBM Accuracy: 0.9504285714285714

In [38]: predictions\_rf = rf\_cv.predict(x\_test)

```

# Classification Report of Prediction
print("Classification Report:")
print(classification_report(y_test, predictions_rf))
# Confusion Matrix for predictions made
conf2 = confusion_matrix(y_test, predictions_rf)
print(conf2)
# Plot Confusion Matrix
label = ["0", "1"]
sns.heatmap(conf2, annot=True, xticklabels=label, yticklabels=label)

```

Classification Report:

	precision	recall	f1-score	support
0	0.95	1.00	0.97	9564
1	0.98	0.79	0.87	2436
micro avg	0.95	0.95	0.95	12000
macro avg	0.96	0.89	0.92	12000
weighted avg	0.95	0.95	0.95	12000


```

[[9521  43]
 [ 519 1917]]

```

Out[38]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1a0af2a2b0>





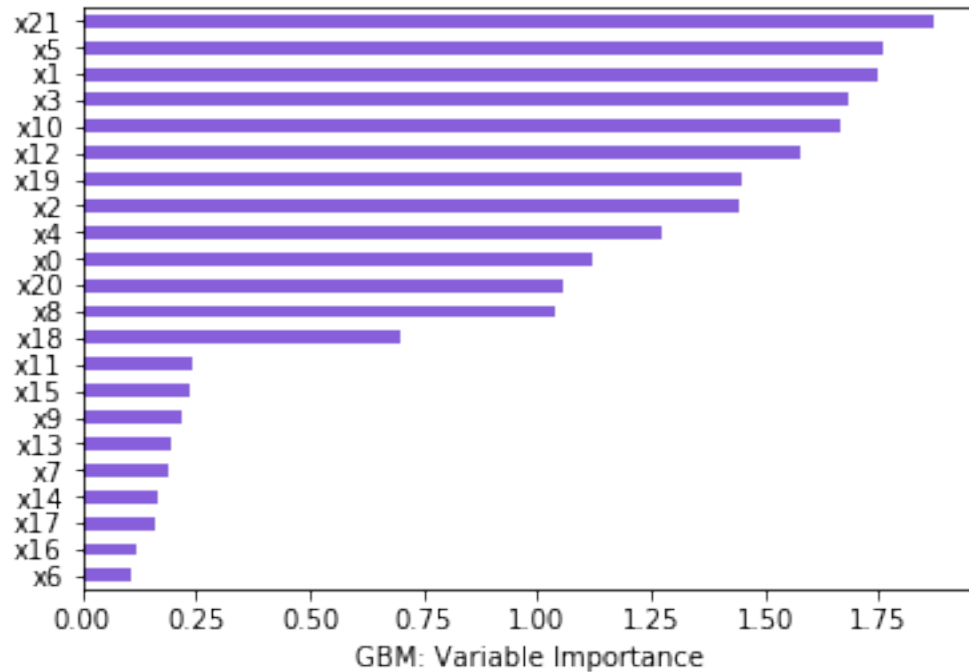
```
In [39]: floatlist = train.loc[:, train.dtypes == float].columns.tolist()
        numcat = rf_cv.best_estimator_.named_steps['preprocessor'].named_transformers_
        bb=numcat['cat'].named_steps["onehot"].get_feature_names()
        feature = np.concatenate((floatlist,bb))
        print(feature.shape)
```

(126,)

```
In [40]: RF = rf_cv.best_estimator_.named_steps['classifier']
```

```
In [41]: plt.figure(figsize=(110,120))
        Importance = pd.DataFrame({'Importance':RF.feature_importances_[0:22]*100},
                                index = feature[0:22])
        Importance.sort_values(
            'Importance', axis=0, ascending=True).plot(kind='barh', color='#875FDB')
        plt.xlabel('GBM: Variable Importance')
        plt.gca().legend_ = None
        plt.savefig('RBM')
```

<Figure size 7920x8640 with 0 Axes>



```
In [42]: testy_rf_pred_proba = rf_cv.predict_proba(testdata)[: ,1]
         pd.DataFrame(testy_rf_pred_proba).to_csv('resultrgbm.csv')
```

```
In [43]: # KS test for training and holdout data
         from scipy import stats
         numeric_features = train.loc[:, train.dtypes == float].columns.tolist()
         trainnum = train[numeric_features]
         testnum = testdata[numeric_features]
         pv = np.empty([len(trainnum.columns), 1])
         t = 0
         for column in trainnum:
             pv[t]=stats.ks_2samp(trainnum[column], testnum[column]).pvalue
             t = t+1
```

## 13 Appendix

### 13.1 Random Forest

I used pipelines to build the Random Forest (RF) classifier. I first created the preprocessing pipelines for both numerical and categorical data. For numerical data, I imputed the missing data with a median and standardized the data. For categorical data, I imputed missing value with the most frequent values and created dummy variable. I used Synthetic Minority Over-sampling (SMOTE) to create new cases make the classification categories equally represented. I build the RF classifier, tuning hyper-parameters with grid search on the number of trees, quality of split criteria, number of features to consider when looking for the best split, with 5-fold cross-validation. The best performing RF uses 500 trees, entropy as split criteria, and the square root of the total number of features as the number of features to consider when looking for the best split. The accuracy score is 0.925, and the ROC AUC score is 0.9752.

In [31]: `from imblearn.pipeline import Pipeline`

```
# We create the preprocessing pipelines for both numeric and categorical data.
```

```
numeric_features = train.loc[:, train.dtypes == float].columns.tolist()
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])
```

```
categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])
```

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])
```

```
# Append classifier to preprocessing pipeline.
```

```
# Now we have a full prediction pipeline.
```

```
os = SMOTENC(categorical_features=[34,35,68,93],random_state=0)
rf = RandomForestClassifier()
clf = Pipeline(steps=[('preprocessor', preprocessor),
    ('sampling', SMOTE()),
    ('classifier', rf)])
```

```
param_grid = {
    'classifier__n_estimators': [200, 500],
    'classifier__max_features': ['sqrt', 'log2'],
    'classifier__criterion': ['gini', 'entropy']
}
```

```
rf_cv = GridSearchCV(clf, param_grid, cv = 5)
```

```
rf_cv.fit(x_train, y_train.values.ravel())
```

```
y_pred_proba = rf_cv.predict_proba(x_test)[:,-1]
```

```
# Evaluate test-set roc_auc_score
```

```
rf_roc_auc = roc_auc_score(y_test, y_pred_proba)
```

```
print("model score: %.3f" % rf_cv.score(x_test, y_test))
```

```
# Print roc_auc_score
```

```
print('ROC AUC score: {:.4f}'.format(rf_roc_auc))
```

```
print("Tuned Random Forest Parameter: {}".format(rf_cv.best_params_))
```

```
print("Tuned Random Forest Accuracy: {}".format(rf_cv.best_score_))
```

model score: 0.925

ROC AUC score: 0.9752

Tuned Random Forest Parameter: {'classifier\_\_criterion': 'entropy', 'classifier\_\_max\_features': 'sqrt', 'classifier\_\_n\_estimators': 500}

Tuned Random Forest Accuracy: 0.9211785714285714

In [32]: predictions\_rf = rf\_cv.predict(x\_test)

```
# Classification Report of Prediction
```

```
print("Classification Report:")
```

```
print(classification_report(y_test, predictions_rf))
```

```
# Confusion Matrix for predictions made
```

```
conf2 = confusion_matrix(y_test, predictions_rf)
```

```
print(conf2)
```

```
# Plot Confusion Matrix
```

```
label = ["0", "1"]
```


```
sns.heatmap(conf2, annot=True, xticklabels=label, yticklabels=label)
```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.99	0.95	9564
1	0.96	0.66	0.78	2436
micro avg	0.93	0.93	0.93	12000
macro avg	0.94	0.83	0.87	12000
weighted avg	0.93	0.93	0.92	12000

```
[[9501  63]  
 [ 836 1600]]
```

Out[32]: <matplotlib.axes.\_subplots.AxesSubplot at 0x10490d518>



```
In [33]: floatlist = train.loc[:, train.dtypes == float].columns.tolist()
numcat = rf_cv.best_estimator_.named_steps['preprocessor'].named_transformers_
bb=numcat['cat'].named_steps["onehot"].get_feature_names()
feature = np.concatenate((floatlist,bb))
print(feature.shape)
feature
```

(126,)

```
Out[33]: array(['x0', 'x1', 'x2', 'x3', 'x4', 'x5', 'x6', 'x7', 'x8', 'x9', 'x10',
                'x11', 'x12', 'x13', 'x14', 'x15', 'x16', 'x17', 'x18', 'x19',
                'x20', 'x21', 'x22', 'x23', 'x24', 'x25', 'x26', 'x27', 'x28',
                'x29', 'x30', 'x31', 'x32', 'x33', 'x36', 'x37', 'x38', 'x39',
                'x40', 'x41', 'x42', 'x43', 'x44', 'x45', 'x46', 'x47', 'x48',
                'x49', 'x50', 'x51', 'x52', 'x53', 'x54', 'x55', 'x56', 'x57',
                'x58', 'x59', 'x60', 'x61', 'x62', 'x63', 'x64', 'x65', 'x66',
                'x67', 'x69', 'x70', 'x71', 'x72', 'x73', 'x74', 'x75', 'x76',
                'x77', 'x78', 'x79', 'x80', 'x81', 'x82', 'x83', 'x84', 'x85',
                'x86', 'x87', 'x88', 'x89', 'x90', 'x91', 'x92', 'x94', 'x95',
                'x96', 'x97', 'x98', 'x99', 'x0_Honda', 'x0_Toyota', 'x0_bmw',
                'x0_chevrolet', 'x0_chrysler', 'x0_ford', 'x0_mercedes',
                'x0_nissan', 'x0_tesla', 'x0_volkswagon', 'x1_friday', 'x1_monday',
                'x1_thursday', 'x1_tuesday', 'x1_wednesday', 'x2_Apr', 'x2_Aug',
                'x2_Dec', 'x2_Feb', 'x2_Jan', 'x2_Jul', 'x2_Jun', 'x2_Mar',
                'x2_May', 'x2_Nov', 'x2_Oct', 'x2_Sep', 'x3_america', 'x3_asia',
                'x3_euorpe'], dtype=object)
```


```

In [34]: RF = rf_cv.best_estimator_.named_steps['classifier']

In [35]: plt.figure(figsize=(110,120))
Importance = pd.DataFrame({'Importance':RF.feature_importances_[0:22]*100},
                           index = feature[0:22])
Importance.sort_values(
    'Importance', axis=0, ascending=True).plot(kind='barh', color='#875FDB')
plt.xlabel('RF: Variable Importance')
plt.gca().legend_ = None
plt.savefig('RF')

```

<Figure size 7920x8640 with 0 Axes>



```

In [36]: testy_rf_pred_proba = rf_cv.predict_proba(testdata)[:,:1]
pd.DataFrame(testy_rf_pred_proba).to_csv('resultrfsmote.csv')

```

## 13.2 K-Nearest Neighbors

I used pipelines to build the *K*-Nearest Neighbors (KNN) classifier. I first created the preprocessing pipelines for both numerical and categorical data. For numerical data, I imputed the missing data with a median and standardized the data. For categorical data, I imputed missing value with the most frequent values and created dummy variable. I build the KNN classifier, tuning hyper-parameters with grid search on the number of neighbors and metric type with 5-fold cross-validation. The best performing KNN uses five neighbors and Euclidean distance. The accuracy

score is 0.9263, and the ROC AUC score is 0.9597. When using SMOTE, the accuracy score is , and the ROC AUC score is .

In [28]: *# Create the preprocessing pipelines for both numeric and categorical data.*

```
numeric_features = floatlist
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler()))

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
knn = KNeighborsClassifier()
clf = Pipeline(steps=[('preprocessor', preprocessor),
    #('sampling', SMOTE()), #use SMOTE or not
    ('classifier', knn)])

param_grid = [{'classifier__n_neighbors': [5,10, 15, 20, 25],
    "classifier__metric": ["euclidean", "cityblock"]
    }]
knn_cv = GridSearchCV(clf, param_grid, cv = 5)

knn_cv.fit(x_train, y_train)

predictions_knn = knn_cv.predict(x_test)
y_pred_proba = knn_cv.predict_proba(x_test)[:,1]
# Evaluate test-set roc_auc_score
knn_roc_auc = roc_auc_score(y_test, y_pred_proba)

# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(knn_roc_auc))
print("Tuned KNN Parameter: {}".format(knn_cv.best_params_))
print("Tuned KNN Accuracy: {}".format(knn_cv.best_score_))
```

ROC AUC score: 0.9611

Tuned KNN Parameter: {'classifier\_\_metric': 'euclidean', 'classifier\_\_n\_neighbors': 5}

Tuned KNN Accuracy: 0.9216785714285715

```

In [29]: predictions_knn = knn_cv.predict(X_test)

# Classification Report of Prediction
print("Classification Report:")
print(classification_report(y_test, predictions_knn))
# Confusion Matrix for predictions made
conf2 = confusion_matrix(y_test, predictions_knn)
print(conf2)
# Plot Confusion Matrix
label = ["0", "1"]
sns.heatmap(conf2, annot=True, xticklabels=label, yticklabels=label)

```

Classification Report:


	precision	recall	f1-score	support
0	0.92	0.99	0.96	9564
1	0.97	0.67	0.79	2436
micro avg	0.93	0.93	0.93	12000
macro avg	0.95	0.83	0.88	12000
weighted avg	0.93	0.93	0.92	12000

```

[[9512  52]
 [ 801 1635]]

```

Out[29]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1131db080>





```
In [30]: testy_logreg_pred_proba = knn_cv.predict_proba(testdata)[: ,1]
pd.DataFrame(testy_logreg_pred_proba).to_csv('resultknnfinal.csv')
```

### 13.3 KNN with PCA

I used pipelines to build the logistic regression. I build KNN after applying PCA and used grid search on hyperparameter number of neighbors for KNN and number of components for PCA with 5-fold cross-validation.

```
In [8]: # Create the preprocessing pipelines for both numeric and categorical data.
numeric_features = floatlist
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
knn = KNeighborsClassifier()
pca = PCA()
clf = Pipeline(steps=[('preprocessor', preprocessor),
                      ('pca', pca),
                      ('classifier', knn)])
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3, stratify = y, random_state = 42)

param_grid = [{'pca__n_components': [ 10, 50, 80,100,110, 124],
                'classifier__n_neighbors': [5, 10, 20]}]
knn_cv = GridSearchCV(clf, param_grid, cv = 2)

knn_cv.fit(X_train, y_train)

predictions_knn = knn_cv.predict(X_test)
y_pred_proba = knn_cv.predict_proba(X_test)[: ,1]
# Evaluate test-set roc_auc_score
knn_roc_auc = roc_auc_score(y_test, y_pred_proba)

# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(knn_roc_auc))
#print("model score: %.3f" % clf.score(X_test, y_test))
```

```

#lg = logreg_cv.best_estimator_.named_steps['classifier']
#print("The Number of Active Features is : {}".format(sum(lg.coef_ == 0).sum()))
print("Tuned Logistic Regression Parameter: {}".format(knn_cv.best_params_))
print("Tuned Logistic Regression Accuracy: {}".format(knn_cv.best_score_))

```

ROC AUC score: 0.9611

Tuned Logistic Regression Parameter: {'classifier\_\_n\_neighbors': 5, 'pca\_\_n\_components': 124}

Tuned Logistic Regression Accuracy: 0.9147857142857143

### 13.4 Logistic Regression with L1

I used pipelines to build the logistic regression. I build logistic regression with L1 regularization and used grid search on hyperparameter C with 5-fold cross-validation.

In [46]: *# Create the preprocessing pipelines for both numeric and categorical data.*

```

numeric_features = train.loc[:, train.dtypes == float].columns.tolist()
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='mean')),
    #('imputer', interpolate(method='linear', inplace=True, limit_direction="both"))
    ('scaler', StandardScaler())])

```

```

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')), #,
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

```

```

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

```

*# Append classifier to preprocessing pipeline.*

*# Now we have a full prediction pipeline.*

```

logreg = LogisticRegression(penalty = 'l1', solver = 'liblinear', max_iter = 500)

```

```

clf = Pipeline(steps=[('preprocessor', preprocessor),
    ('classifier', logreg)])

```

```

X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
    stratify = y, random_state = 42)

```

```

#clf.fit(X_train, y_train)

```

```

c_space = np.logspace(-5, 8, 15)

```

```

param_grid = {'classifier__C': c_space} #{'classifier__C': c_space}

```

```

logreg_cv = GridSearchCV(clf, param_grid, cv = 5)

```

```

logreg_cv.fit(X_train, y_train)

```

```

y_pred_proba = logreg_cv.predict_proba(X_test)[:,:1]

```

```

# Evaluate test-set roc_auc_score
logreg_roc_auc = roc_auc_score(y_test, y_pred_proba)

# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(logreg_roc_auc))
print("model score: %.3f" % logreg_cv.score(X_test, y_test))
lg = logreg_cv.best_estimator_.named_steps['classifier']
print("The Number of Active Features is : {}".format(sum(lg.coef_ != 0).sum()))
print("Tuned Logistic Regression Parameter: {}".format(logreg_cv.best_params_))
print("Tuned Logistic Regression Accuracy: {}".format(logreg_cv.best_score_))

```

ROC AUC score: 0.9088

model score: 0.891

The Number of Active Features is : 106

Tuned Logistic Regression Parameter: {'classifier\_\_C': 0.4393970560760795}

Tuned Logistic Regression Accuracy: 0.8904285714285715

### 13.5 Logistic Regression with Recursive Feature Elimination

I used pipelines to build the logistic regression. I first created the preprocessing pipelines for both numerical and categorical data. For numerical data, I imputed the missing data with median and standardized the data. For categorical data, I imputed missing value as missing and created dummy variable (ignoring the missing data). I split the data into 70% training data and 30% test data. I build logistic regression with Recursive Feature Elimination (RFE) and used grid search on hyperparameter C for logistic regression and number of features for RFE with 5-fold cross-validation.

```

In [ ]: # Create the preprocessing pipelines for both numeric and categorical data.
numeric_features = train.loc[:, train.dtypes == float].columns.tolist()
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
logreg = LogisticRegression(solver='lbfgs', max_iter=300)
rfe = RFE(logreg)

```

```

clf = Pipeline(steps=[('preprocessor', preprocessor),
                      ('rfe', rfe),
                      ('classifier', logreg)
                      ])
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
                                                  stratify = y, random_state = 42)

c_space = np.logspace(-5, 8, 15)
param_grid = {'classifier__C': [.01, .5],
              'rfe__n_features_to_select': [30, 50]}
logreg_cv = GridSearchCV(clf, param_grid, cv = 5)

logreg_cv.fit(X_train, y_train)

y_pred_proba = logreg_cv.predict_proba(X_test)[:,-1]
# Evaluate test-set roc_auc_score
logreg_roc_auc = roc_auc_score(y_test, y_pred_proba)

# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(logreg_roc_auc))
#print("model score: %.3f" % logreg_cv.score(X_test, y_test))
lg = logreg_cv.best_estimator_.named_steps['classifier']
print("The Number of Active Features is : {}".format(sum(lg.coef_ != 0).sum()))
print("Tuned Logistic Regression Parameter: {}".format(logreg_cv.best_params_))
print("Tuned Logistic Regression Accuracy: {}".format(logreg_cv.best_score_))

```

ROC AUC score: 0.9069

The Number of Active Features is : 50

Tuned Logistic Regression Parameter: {'classifier\_\_C': 0.5, 'rfe\_\_n\_features\_to\_select': 50}

Tuned Logistic Regression Accuracy: 0.8896832773289484

## 13.6 PCA + Logistic Regression

I first applied PCA to the data and build logistic regression on the transformed features from PCA, and used grid search on hyperparameter C for logistic regression and number of components for PCA with 5-fold cross-validation. Uncomment the SMOTE commands in the pipeline will use SMOTE in the analysis.

In [11]: *# Create the preprocessing pipelines for both numeric and categorical data.*

```

numeric_features = floatlist
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])

categorical_features = ['x34', 'x35', 'x68', 'x93']
categorical_transformer = Pipeline(steps=[

```

```

('imputer', SimpleImputer(strategy='most_frequent')),
('onehot', OneHotEncoder(handle_unknown='ignore'))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Append classifier to preprocessing pipeline.
# Now we have a full prediction pipeline.
logistic = SGDClassifier(loss='log', penalty='l2', early_stopping=True,
                        max_iter=10000, tol=1e-5, random_state=0)

pca = PCA()
clf = Pipeline(steps=[('preprocessor', preprocessor),
                      #('sampling', SMOTE()), #use SMOTE or not
                      ('pca', pca),
                      ('classifier', logistic)])
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.3,
                                                    stratify = y, random_state = 42)

param_grid = [{'pca__n_components': [50, 80, 90, 100],
               'classifier__alpha': np.logspace(-4, 4, 5)}] #n
knn_cv = GridSearchCV(clf, param_grid, cv = 2)

knn_cv.fit(X_train, y_train)

predictions_knn = knn_cv.predict(X_test)
y_pred_proba = knn_cv.predict_proba(X_test)[:,-1]
# Evaluate test-set roc_auc_score
knn_roc_auc = roc_auc_score(y_test, y_pred_proba)

# Print roc_auc_score
print('ROC AUC score: {:.4f}'.format(knn_roc_auc))
#print("model score: %.3f" % clf.score(X_test, y_test))
#lg = logreg_cv.best_estimator_.named_steps['classifier']
#print("The Number of Active Features is : {}".format(sum(lg.coef_ == 0).sum()))
print("Tuned Logistic Regression Parameter: {}".format(knn_cv.best_params_))
print("Tuned Logistic Regression Accuracy: {}".format(knn_cv.best_score_))

```

ROC AUC score: 0.9036

Tuned Logistic Regression Parameter: {'classifier\_\_alpha': 0.01, 'pca\_\_n\_components': 100}

Tuned Logistic Regression Accuracy: 0.8848214285714285